

Within this project, I would like to explore the CoinGecko API and the package “geckor” in R. Then I would collect the current and historical cryptocurrency market data for a coin using the public ‘CoinGecko’ API and do a Arima forecast.

Let’s use bitcoin as an example here. First, I want to obtain historical data for bitcoin. The function “coin_history” can retrieve coin-specific market data for the last n days. If open-high-low-close price data is needed, use function “coin_history_ohlcv” instead.

```
r <- coin_history(coin_id = "bitcoin", vs_currency = "usd", days = "max")
```

```
## Warning: Missing values found in column(s)
## * market_cap
```

```
r
```

```
## # A tibble: 3,435 x 6
##   timestamp          coin_id vs_currency price total_volume market_cap
##   <dtm>              <chr>   <chr>    <dbl>      <dbl>      <dbl>
## 1 2013-04-28 00:00:00 bitcoin usd      135.        0 1500517590
## 2 2013-04-29 00:00:00 bitcoin usd      142.        0 1575032004
## 3 2013-04-30 00:00:00 bitcoin usd      135.        0 1501657493
## 4 2013-05-01 00:00:00 bitcoin usd      117.        0 1298951550
## 5 2013-05-02 00:00:00 bitcoin usd      103.        0 1148667722
## 6 2013-05-03 00:00:00 bitcoin usd       91.0        0 1011066494
## 7 2013-05-04 00:00:00 bitcoin usd      111.        0 1236351844
## 8 2013-05-05 00:00:00 bitcoin usd      117.        0 1298377788
## 9 2013-05-06 00:00:00 bitcoin usd      118.        0 1315992304
## 10 2013-05-07 00:00:00 bitcoin usd      106.        0 1183766500
## # ... with 3,425 more rows
```

Since we only need timestamp and price from obtained data, I save them in a dataframe and convert it to timeserie format for later use. So now df has all daily market price for bitcoin. Below is last 7 days price in ts object.

```
var <- c("timestamp", "price")
df <- r[var]
df <- df[c(1:nrow(df) - 1),]
date <- df$timestamp
start_date <- date[1]
dayOfYear <- as.numeric(format(as.Date(start_date), "%j"))
year <- as.numeric(format(as.Date(start_date), "%Y"))
df <- ts(df$price, start = c(year, dayOfYear), frequency = 365)
tail(df, 7)
```

```
## Time Series:
## Start = c(2022, 260)
## End = c(2022, 266)
## Frequency = 365
## [1] 19764.41 20131.68 19437.16 19570.39 18869.93 18539.64 19464.32
```

Now, let’s define the training and testing period. I will use the last 7 days as testing set and all other historical data as training set.

```
n = length(df)
n1 = n - 6
n2 = n - 7

df_train = df[c(1:n2)]
df_test = df[c(n1:n)]
```

We now need to check the stationarity of our data. We can do that with the Augmented Dickey-Fuller Test.

H0: The time series is non-stationary. H1: The time series is stationary. Since the p-value is not less than .05, we fail to reject the null hypothesis. This means the time series is non-stationary. Our data is dependent on the time at which the series is observed.

```
adf.test(df)

##
## Augmented Dickey-Fuller Test
##
## data: df
## Dickey-Fuller = -2.1268, Lag order = 15, p-value = 0.5246
## alternative hypothesis: stationary
```

We will run the `auto.arima` function on our training data, which will help us to forecast next 7 days market price and their prediction intervals.

```
model <- auto.arima(df_train)
fcast <- forecast(model, h = 7, level = 95)

fcast
```

```
##      Point Forecast      Lo 95      Hi 95
## 3428      19702.17 18241.79 21162.55
## 3429      19702.17 17636.88 21767.46
## 3430      19702.17 17172.71 22231.63
## 3431      19702.17 16781.41 22622.93
## 3432      19702.17 16436.66 22967.68
## 3433      19702.17 16124.98 23279.36
## 3434      19702.17 15838.36 23565.98
```

let's create a table of predicted price and actual price.

```
result <- cbind(df_test, as.numeric(fcast$mean))
colnames(result) <- c("Actual", "Predicted")
result
```

```
##      Actual Predicted
## [1,] 19764.41 19702.17
## [2,] 20131.68 19702.17
## [3,] 19437.16 19702.17
## [4,] 19570.39 19702.17
## [5,] 18869.93 19702.17
## [6,] 18539.64 19702.17
## [7,] 19464.32 19702.17
```

Now, let's get the Metrics MAPE which is used to judge the performance of the model. It gives the average deviation between the forecast value and actual values.

```
mape <- mape(df_test, as.numeric(fcast$mean))  
mape*100
```

```
## [1] 2.34116
```

Lastly, plot the historical price for bitcoin, certainly the price has been jumping insane last few years.

```
autoplot(fcast) + (labs(y = "Price", x = "Days"))
```

