

Within this project, I would like to explore the CoinGecko API and the package “geckor” in R. Then I would collect the current and historical cryptocurrency market data for a coin using the public ‘CoinGecko’ API and do a Arima forecast.

## Data preparation

Let’s use dogecoin as an example here. First, I want to obtain historical data for dogecoin. The function “coin\_history” can retrieve coin-specific market data for the last n days. If open-high-low-close price data is needed, use function “coin\_history\_ohlcv” instead.

```
r <- coin_history(coin_id = "dogecoin", vs_currency = "usd", days = "max")
```

```
## Warning: Missing values found in column(s)
## * market_cap
```

```
r
```

```
## # A tibble: 3,204 x 6
##   timestamp          coin_id vs_currency   price total_volume market_cap
##   <dtm>              <chr>   <chr>      <dbl>      <dbl>      <dbl>
## 1 2013-12-15 00:00:00 dogecoin usd      0.000559         0    3488670
## 2 2013-12-17 00:00:00 dogecoin usd      0.000218         0    1619159
## 3 2013-12-18 00:00:00 dogecoin usd      0.000268         0    2191987
## 4 2013-12-19 00:00:00 dogecoin usd      0.000475         0    4299422
## 5 2013-12-20 00:00:00 dogecoin usd      0.000989         0    9866232
## 6 2013-12-21 00:00:00 dogecoin usd      0.000438         0    4686300
## 7 2013-12-22 00:00:00 dogecoin usd      0.000405         0    4639566
## 8 2013-12-23 00:00:00 dogecoin usd      0.000330         0    4022744
## 9 2013-12-24 00:00:00 dogecoin usd      0.000712         0    9213651
## 10 2013-12-25 00:00:00 dogecoin usd      0.000582         0    7955874
## # ... with 3,194 more rows
```

Since we only need date and price from obtained data, I save them in a dataframe and convert it to timeserie format for later use. So now df has all daily market price for dogecoin. Below is last 7 days price in ts object. Since each coin has different start date in coin market, we need to record the first date in df.

```
var <- c("timestamp", "price")
df <- r[var]
df <- df[c(1:nrow(df) - 1),]
date <- df$timestamp
start_date <- date[1]
dayOfYear <- as.numeric(format(as.Date(start_date), "%j"))
year <- as.numeric(format(as.Date(start_date), "%Y"))
df <- ts(df$price, start = c(year, dayOfYear), frequency = 365)
tail(df, 7)
```

```
## Time Series:
## Start = c(2022, 260)
## End = c(2022, 266)
## Frequency = 365
## [1] 0.06229599 0.05760080 0.05871669 0.05851463 0.05749412 0.05979688 0.06337732
```

Now, let's define the training and test period. I will use the last 7 days as test set and all other historical data as training set.

```
train_end <- length(df) - 7
test_start <- length(df) - 6

df_train <- ts(df[c(1:train_end)])
df_test <- ts(df[c(test_start:length(df))])
```

We now need to check the stationarity of our data. We can do that with the Augmented Dickey-Fuller Test.

H0: The time series is non-stationary. H1: The time series is stationary. Since the p-value is not less than .05, we fail to reject the null hypothesis. This means the time series is non-stationary. Our data is dependent on the time at which the series is observed.

```
adf.test(df)

##
## Augmented Dickey-Fuller Test
##
## data: df
## Dickey-Fuller = -2.7914, Lag order = 14, p-value = 0.2432
## alternative hypothesis: stationary
```

## Time series modeling

### ARIMA models which aim to describe the autocorrelations in the data.

We will run the `auto.arima` function on our training data, which will help us to forecast next 7 days market price and their prediction intervals.

```
fit_arima <- auto.arima(df_train)
fcast_arima <- forecast(fit_arima, h = 7, level = 95)

fcast_arima
```

```
##      Point Forecast      Lo 95      Hi 95
## 3197      0.06084483 0.04389762 0.07779204
## 3198      0.06193122 0.03977521 0.08408722
## 3199      0.06277872 0.03536605 0.09019138
## 3200      0.06299779 0.03053659 0.09545898
## 3201      0.06253890 0.02568147 0.09939633
## 3202      0.06168589 0.02133391 0.10203786
## 3203      0.06088812 0.01787553 0.10390072
```

Now, let's get the Metrics MAPE which is used to judge the performance of the model. It gives the average deviation between the forecast value and actual values.

```
mape_arima <- mape(df_test, as.numeric(fcast_arima$mean))*100
mape_arima
```

```
## [1] 5.75545
```

Exponential smoothing models (ETS) are based on a description of the trend and seasonality in the data.

```
fit_ets <- ets(df_train)
fcast_ets <- forecast(fit_ets, h = 7, level = 95)

mape_ets <- mape(df_test, as.numeric(fcast_ets$mean))*100
mape_ets
```

```
## [1] 3.208157
```

let's create a table of predicted price, actual price and their MAPE at last row. This will help us to do comparison.

```
result <- cbind(df_test, as.numeric(fcast_arima$mean), as.numeric(fcast_ets$mean))
result <- rbind(result, c(0, mape_arima, mape_ets))
colnames(result) <- c("Actual", "ARIMA", "ETS")
round(result, 4)
```

```
##      Actual  ARIMA   ETS
## [1,] 0.0623 0.0608 0.0602
## [2,] 0.0576 0.0619 0.0600
## [3,] 0.0587 0.0628 0.0598
## [4,] 0.0585 0.0630 0.0595
## [5,] 0.0575 0.0625 0.0593
## [6,] 0.0598 0.0617 0.0591
## [7,] 0.0634 0.0609 0.0589
## [8,] 0.0000 5.7554 3.2082
```

Lastly, plot the historical price for bitcoin, certainly the price has been jumping insane last few years.

```
autoplot(fcast_arima) + (labs(y = "Price", x = "Days"))
```

Forecasts from ARIMA(2,1,3)

